



**Department of Electrical,
Computer, & Biomedical Engineering**
Faculty of Engineering
& Architectural Science

Course Number	COE758
Course Title	Digital System Engineering
Semester/Year	F2024
Instructor	Dr. Lev Kirischian
Section No.	01
Group No.	NA
Submission Date	2024/10/15
Due Date	2024/10/15

Assignment Title	Cache Controller Project Report
-------------------------	---------------------------------

Name	Student ID	Signature*
Vaishali Jadon	501051835	V.J.
Atiya Azeez	501052595	A.A.

**By signing above you attest that you have contributed to this submission and confirm that all work you have contributed to this submission is your own work. Any suspicion of copying or plagiarism in this work will result in an investigation of Academic Misconduct and may result in a “0” on the work, an “F” in the course, or possibly more severe penalties, as well as a Disciplinary Notice on your academic record under the Student Code of Academic Conduct, which can be found online at: www.rverson.ca/senate/current/pol60.pdf*

TABLE OF CONTENTS

Abstract	2
Introduction	2
Purpose	3
Theory	4
Internal Specifications	4
Device Design Description	9
Timing Results	12
Qualitative Results	16
Conclusion	16
References	17
Appendix	18

Abstract

Objective

The main goals of this project were to understand how a cache controller works and how it can access data from the SRAM memory (used for the cache) and SDRAM controllers (main memory). It focused on designing custom logic controllers and interfacing them with the SRAM and other logic devices. The project aimed to build skills in VHDL coding using the Xilinx ISE CAD environment, with the Xilinx Spartan-3E FPGA as the hardware platform. This experience helped us learn about memory management and controller design in digital systems.

Approach

The Cache Controller was developed using VHDL. The cache project implemented the CPU's VHDL implementation provided as well as the ICON, ILA, and BRAM cores taught in previous tutorials. In the initial stages, a schematic diagram, block diagram, and conceptual finite state machine (FSM) were created. Following this, the Cache Controller symbol was generated and the VHDL code for it was written. The hardware emulation was observed using the Chipscore Analyzer which displayed the different cases that the cache supports.

Introduction

Modern computer systems are designed with a variety of memory types that are organized in a memory hierarchy to optimize performance and cost efficiency. In such hierarchies, as the distance from the CPU increases, both the capacity and the access time of the memory increase. This structure enables processors to effectively utilize slower, larger, and less expensive memory. This reduces the delays associated with accessing large, slow memories by using a faster, smaller memory close to the processor.

At the top of this hierarchy is the cache memory, a small, high-speed storage area used to hold data that has been accessed recently and is likely to be needed again soon. Data in the cache is stored in blocks, each containing multiple words. To manage these blocks and ensure efficient data access, a cache controller is required. It maintains necessary information, including the block's index, tag, valid bit, and dirty bit, which helps the controller track the contents of the cache and how they map to the main memory. Since there are multiple words in a block, an offset is used to access a specific word.

The cache controller handles read and write requests from the CPU by either retrieving or updating data within the cache or transferring whole blocks of data to and from the main memory as needed. This report details the implementation of a cache controller in VHDL, designed to interface with a pre-existing CPU model. The report covers the development process, including

design diagrams, VHDL coding, hardware emulation results, and analysis of the cache's performance.

Purpose

The purpose of this report is to understand how the concept of memory hierarchies improves the performance and efficiency of computer systems. Memory hierarchies help bridge the gap between the fast processing power of CPUs and the slower speed of main memory. By organizing memory into levels, from the small and fast cache memory close to the CPU to the larger and slower main memory, computer systems can work more efficiently. This report aims to show how good memory hierarchy design can reduce delays in accessing data and make data transfer faster. The goal is to gain a better understanding of how memory management is critical for improving performance.

Theory

Cache memory stores data in blocks, each containing multiple words. These blocks are identified by several key components: the index, tag, valid bit, dirty bit, and offset. The index helps the system locate which block in the cache corresponds to the requested data, while the tag is used to verify that the block contains the correct data. The valid bit indicates whether the block contains valid data, and the dirty bit shows whether the data in the block has been modified and needs to be written back to the main memory. The offset is used to access individual words within the block.

The Cache Controller is responsible for managing data transfers between the CPU, cache memory, and main memory. It handles both read and write requests from the CPU by determining whether the requested data is in the cache (a cache hit) or not (a cache miss). If a cache hit occurs, the controller directly accesses or updates the data in the cache. If a cache miss occurs, the controller fetches the appropriate data block from the main memory or writes modified blocks back to the main memory before updating the cache.

For this project, the cache can store 256 bytes of data. It is structured into 8 blocks, each holding 32 one-byte words. The CPU communicates with the cache controller by sending 16-bit addresses, which are divided into the tag, index, and offset. The Cache Controller uses this address information to compare the tags of the cache blocks, check the valid and dirty bits, and determine the appropriate action for each read or write request.

The behavioral operation of the Cache Controller can be divided into four main cases:

1. **Write a Word to Cache (Cache Hit):** When the CPU issues a write request for data already stored in the cache, it is categorized as a cache hit. The Cache Controller uses the index and offsets from the CPU's address to locate the correct position in the SRAM. The new data is then written to this location, and both the dirty and valid bits for the corresponding block are set to 1. This indicates that the data has been modified and the cache block has been used.
2. **Read a Word from Cache (Cache Hit):** If the CPU requests to read data that exists in the cache, the Cache Controller identifies it as a cache hit. It retrieves the data using the index and offset from the CPU's address and sends the requested data back to the CPU.
3. **Read/Write from/to Cache (Cache Miss, Dirty Bit = 0):** When a read or write request is made but the corresponding block is not found in the cache (the tags do not match), a cache miss occurs. If the dirty bit is 0, the data in the cache has not been modified, thus the Cache Controller can replace the old block with the new block from the main memory. The CPU address is sent with the initial offset set to "00000" to read the entire 32-byte block. This block is then written to the SRAM, and the tag register is updated with the new tag from the CPU address. The valid bit is set to 1, and the procedure for a hit operation is performed.
4. **Read/Write from/to Cache (Cache Miss, Dirty Bit = 1):** When a cache miss occurs and the dirty bit is set to 1, the Cache Controller must first write the modified data back to the main memory. The data block from the SRAM is written to the following address: tag & index & 00000. After writing back, the Cache Controller retrieves the new block from the main memory using the same address format. The cache is then updated with the new data block and the original CPU request is executed.

Internal Specifications

1. CPU Interface

The CPU requests write and read transactions to the cache. The CPU itself has a chip select strobe (CS), a read/write enable (WR/RD), a 16-bit address (ADD), 8-bit data input and output ports: DIN and DOUT, and a ready indicator (RDY). The clock signal (CLK) is synchronized with the Cache Controller and sends read or write instructions to it. The ready signal (RDY) is used to indicate whether the cache is in an idle or busy state. The ports are shown in Figure 1. The CPU interface has an 8-bit data input pin (DIN) that gets data from the Cache (SRAM). It also sends an address to the Controller, along with a read/write signal and a CS strobe. The CS strobe tells the SRAM that the address, read/write signal, and data output are ready to be used for the instruction. The CS signal

stays asserted for four clock cycles. The RDY signal is used by the Cache Controller to communicate the completion of a transaction. When the Cache Controller is idle and ready to accept new transactions, it asserts the RDY signal. When a transaction is received, the RDY signal is deasserted and remains low until the requested operation is fully completed. An example of a write operation is shown in Figure 2.




Figure 1. CPU symbol interface.

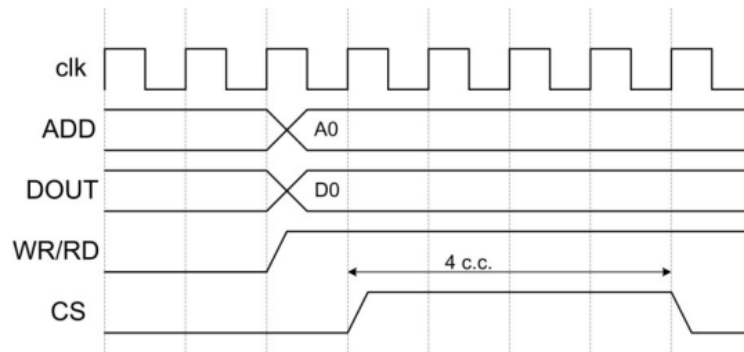


Figure 2. CPU transaction example.

2. SDRAM Controller

The SDRAM Controller processes main memory block read and write requests from the cache. Its block symbol is shown in Figure 7. It consists of a 16-bit address (ADD), a read/write signal (WR/RD), a strobe signal (MEMSTRB), and data input and output ports (DIN and DOUT). The SDRAM Controller and Cache Controller are synchronized using the same clock.

To write a word to the main memory, the Cache Controller sends the block address (tag & index & offset) and sets the WR/RD signal to 1 to indicate a write operation. After ensuring all signals are stable for one clock cycle, the MEMSTRB signal is asserted for

that cycle. Writing a complete block of data to memory requires repeating this process 32 times, with incrementing the offset of the address and adjusting the data accordingly.

When reading a full block from memory, the procedure is similar, except the WR/RD signal is cleared to 0 so no data gets written. Instead, the requested data will be output on the DOUT port. Like writing, reading an entire block takes 32 clock cycles to complete. Figures 8 and 9 below demonstrate the read and write processes for an entire block of data.




Figure 3. SDRAM controller interface.

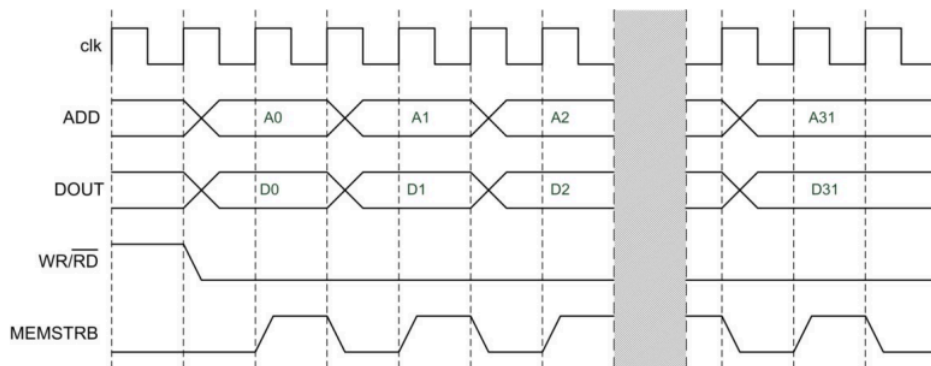


Figure 4. SDRAM block read.

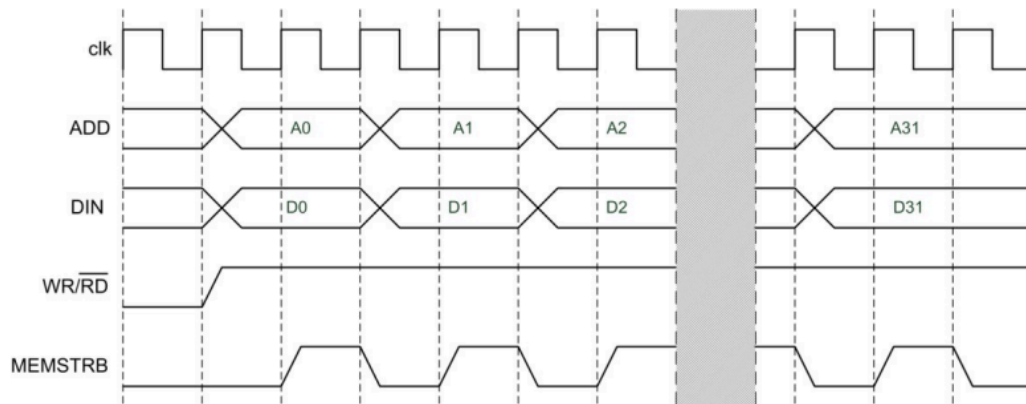


Figure 5. SDRAM block write.

3. Local SRAM

The BlockRAM memory (Figure 6) is used for local cache storage. This interface consists of an 8-bit address (**ADD**), 8-bit data input and output lines (**DIN** and **DOUT**), and a write enable signal (**WEN**). Like the other components, the BlockRAM operates using the same clock signal as the controller. All commands sent to the cache controller are synchronized with the clock's rising edge. For read operations, the correct address is placed on the address bus, and the corresponding data will be available on the **DOUT** line after the next clock-rising edge (Figure 7). Write operations involve specifying both the address and input data, followed by activating the write enable signal (**WEN**). The data will then be written to the designated address during the next rising edge of the clock (Figure 8).




Figure 6. BlockRAM interface




Figure 7. BlockRAM read operation




Figure 8. BlockRAM write operation

Device Design Description

Cache symbol:

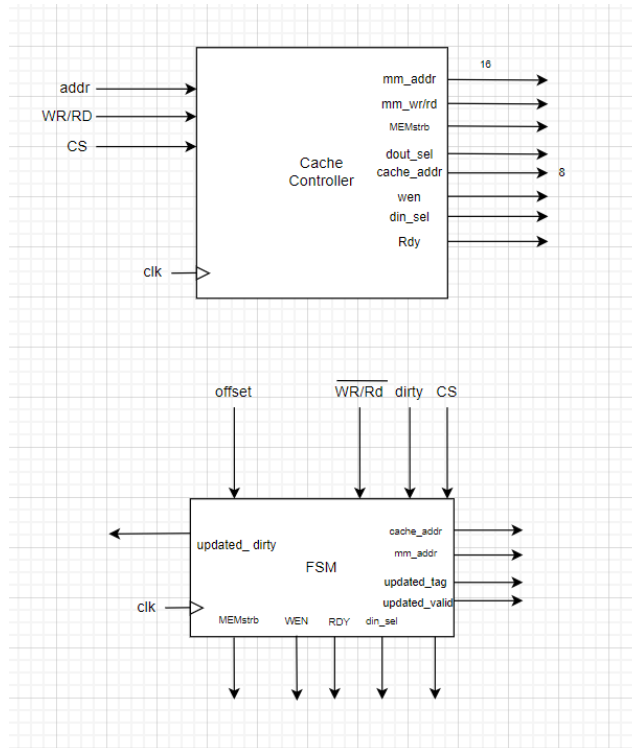


Figure 9. Cache controller

Block diagram:

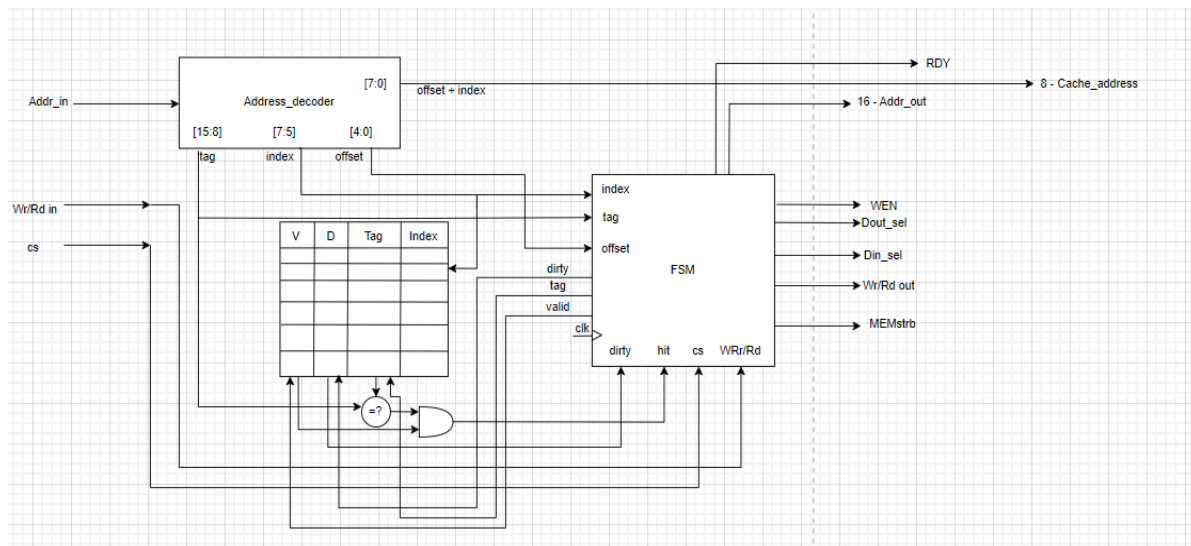


Figure 9. Block diagram of the cache controller

The block diagram of the Cache Controller in Figure 9 illustrates the component's internal workings. The 16-bit Input Address is connected to both the Address Decoder and FSM. The Address Decoder processes the incoming address by breaking it down into three fields: the tag field, the index field, and the offset field. The tag field identifies the block in the cache. The index field is used to locate the corresponding cache block within the cache memory that may contain the requested data, while the offset identifies a specific location within that block. Additionally, there are two values: dirty and valid used to check the data. The valid bit indicates whether data is stored in the cache entry and the dirty bit signifies whether the byte has been modified since it was loaded into memory. The valid bit value is combined with the results of comparing the address and tag assigned to that cache index in a logical AND gate. The outcome of this AND operation is then sent as the “hit” of the FSM. Subsequently, the FSM control generates several output signals: RDY, Addr_out, WEN, Din_sel, Dout_sel, WR/RD, and MEMSTRB.

State Machine Diagram

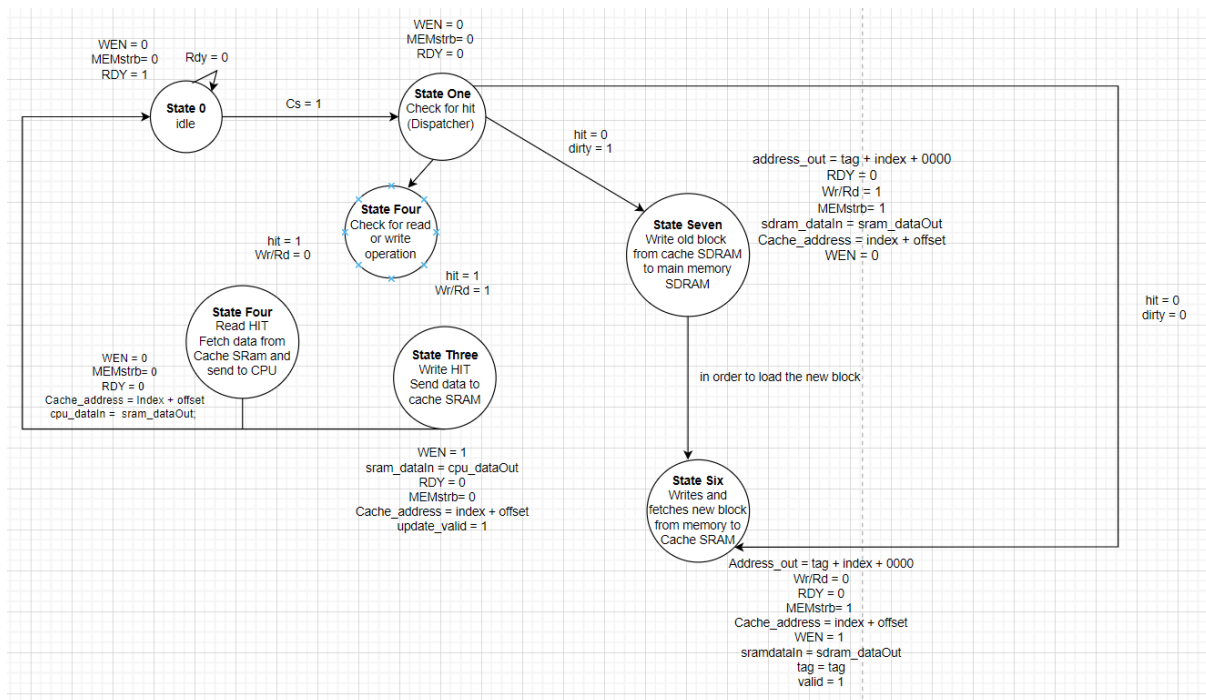


Figure 10. Finite state machine

The FSM in Figure 10, outlines the different states and transitions used by the controller. It consists of seven states, each accompanied by specific conditions for transitions and control signals. The states are outlined as follows:

State 0: Check if the cache is idle. If the CPU sends a request to the cache (CS=1), then the cache can accept this request. The ready signal is deasserted once the cache starts to process the CPU request.

State 1: Decodes the address sent from the cpu into its respective tag, index, and offset values and checks for a hit or miss by comparing the tags and valid bits. The cache tag is taken from the tag table and the valid bit is taken from the valid signal. The WEN signal is set to zero, to ensure that no data is written to cache before the checking occurs. The MEMSTRB signal is deasserted since no data is being read or written from the main memory.

State 2: If a hit occurs, this state determines whether it is a read or write operation that is performed. It uses the read and write enable signal (wr/rd) to determine this.

State 3: Performs the write operation, writing the data provided from the CPU to the specified cache address. Since a write is performed, the WEN signal needs to be asserted. In addition, since the cache data is updated, the valid is set to one for that cache block. Since the CPU request is completed, the cache returns to an idle state (state zero) making it ready to accept other requests.

State 4: Performs the read operation from the address that the CPU requested from the cache. In this case, the data out from the SRAM is outputted to the CPU Data In port. Since the CPU request is completed, the cache returns to an idle state ready to accept other requests.

State 7: If there is a miss and a dirty bit is one, then the cache block gets loaded into the memory. Since the SDRAM is written to, the SDRAM wr/rd enable signal must be asserted. To ensure that data is only written on the rising edge, the data is only read when the MEMSTRB is set to one.

State 6: If the dirty bit is zero, then the new memory block that the CPU address is included in, gets written to the cache. For this operation, the SDRAM wr/rd enable signal is deasserted since only data is being read. Once again, to ensure data output is synchronized, the SDRAM will only output data on the rising edge, when MEMSTRB is one. Moreover, the WEN for SRAM is set to one since the new block is being written and the valid bit is set to 1 to indicate that this cache block has data. Once this is completed, this state goes to state 2 to determine whether a read or write is performed.

Timing Results

1. Cache read operation

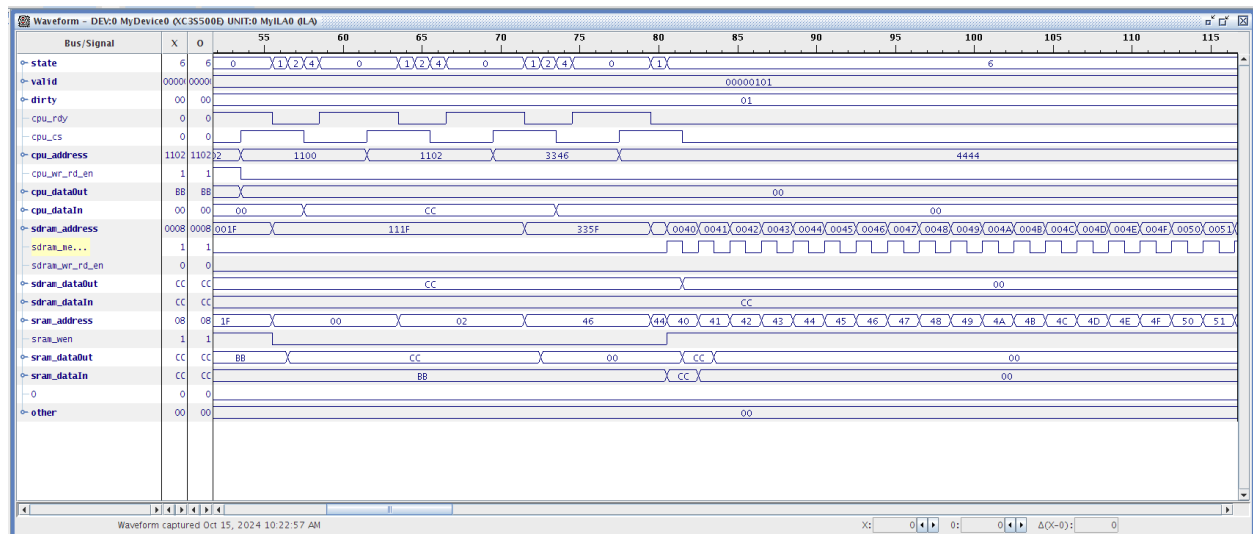


Figure 11. Cache operation (state 4) waveform

The waveform in Figure 11 shows a cache read operation (x=60). The waveform starts from state 0, which is the idle state. Once `cpu_cs` is set to one, the FSM proceeds to state 1 where `cpu_rdy` is 0 because the cache is not available. In state 1, the cache controller checks whether the requested address from the CPU results in a cache hit or miss. The `cpu_address` is split into the tag, index, and offset: 0001 001, 000, and 00010. Hence, the `sram_address` is 00000010 (0x02). This time around, it's a hit, because the valid bit is 1, and the tag matches. Since a cache hit occurs, the FSM moves to state 2. In this state, the controller checks if the CPU is requesting to read or write and this will be based on the `cpu_wr_rd_en` signal. From this waveform, since `cpu_wr_rd_en` is 0, the FSM moves onto state 4 which will handle reading data from the cache (SRAM) and sending that information to the CPU. The `cpu_dataIn` signal is loaded with the value from `sram_dataout`, in this case, it is 0xCC. Since no writes are required, `sram_we` and `sdras_wr_rd_en` are 0. After finishing the read operation, the FSM returns to state 0, reasserting `cpu_ready` as it waits for the next request.

2. Cache write operation




Figure 12. Cache operation (state 3) waveform

The following waveform shows the write operation (starting from x=148). Like the read operation, the waveform goes from state 0, which is the idle state, and then proceeds to state 1. In state 1, the cache controller checks whether the requested address from the CPU results in a cache hit or miss using the tag and index from the `cpu_address`: 0101 0101, 000, 00100. In this case, it is a hit, because the valid bit is 1 and the tag matches. This means that the FSM moves to state 2 to determine the read or write operation. Since `cpu_wr_rd_en` is 1, the CPU is requesting a write operation. Now we move to state 3 where the CPU is writing into the cache. From the CPU address fields, the `sram_address` being written to is 00100 (0x04). The controller will enable write to cache by setting `sram_wen` to 1. The CPU's data (`cpu_dataout`) will be written to the cache block (`sram_dataIn`). In this case, the value 0xCC is being written. It will update the valid bit to 1 and this will indicate that the cache block contains valid data. The dirty bit is also set to 1 indicating that the cache block has been changed. After state 3, the FSM moves to state 0 and waits for the next request

3. Dirty Bit is 0

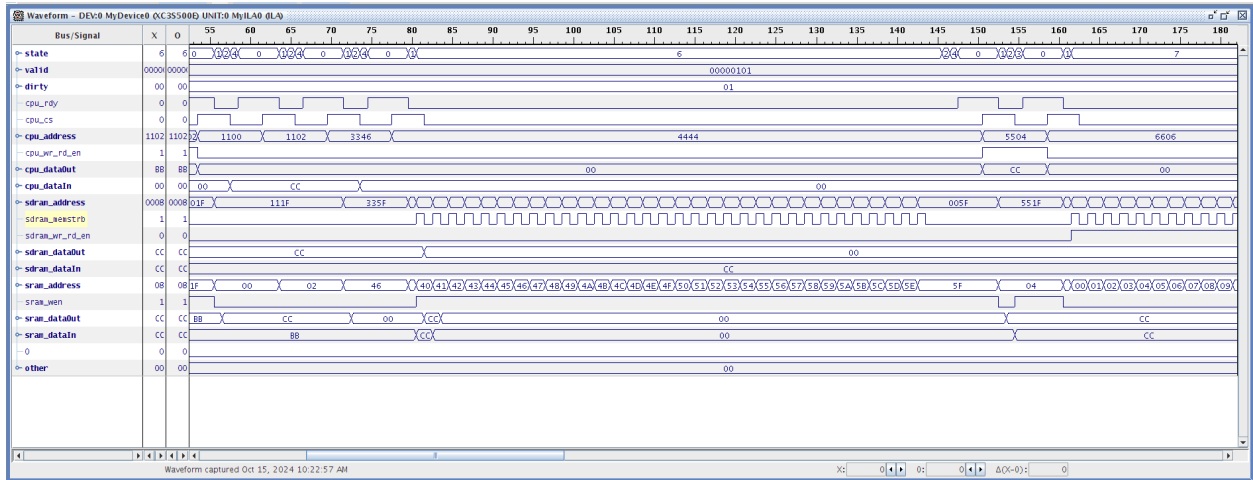


Figure 13. Cache operation (state 6) waveform

The waveform above is when the dirty bit is 0. At $x = 75$, the FSM is at state 0 and the cache is waiting for the CPU to send a signal command (`cpu_cs`). Once the signal is received, the FSM asserts `cpu_rdy` to 1 and that indicates the cache is performing the request. In state 1, the `cpu_address` is loaded and the FSM extracts the tag, index, and offset from the address: 0100 0100, 010, and 000100. The FSM will also check if the requested data is valid in the cache and whether the block is dirty. In this waveform, the valid bit is 0 (the second bit of valid) and the dirty bit is 0 (the second bit of diary). Thus, it is a cache miss. Since the dirty bit is 0, it proceeds to state 6 to load the missing block from the main memory (SDRAM) into the cache (SRAM). A 64-cycle counter is used to track the 32-byte data transfers since a full block write takes 32 clock cycles (one block per cycle). The cache controller loads a new block from SDRAM to SRAM, byte by byte over multiple clock cycles. This is seen as the SRAM address increments by one each cycle as it writes each byte from memory. In addition, the waveform shows that data outputs from the SDRAM when the `sdran_memstrb` is on the rising edge. Since the SRAM is being written to, `sram_wen` is set to one and since SDRAM is being read from, `sdran_wr_rden` is set to 0. Once the counter hits 64, the transfer is complete and the FSM transitions to state 2. Now at state 2, it will handle the CPU's original request. FSM will either read or write operations depending on the value of the `cpu_wrd_en` signal. Since `cpu_wrd_en` = 0, the FSM will move to state 4 read data from the cache, and move it forward to the CPU. The FSM will then return to state 0 in its idle state.

4. Dirty bit is 1

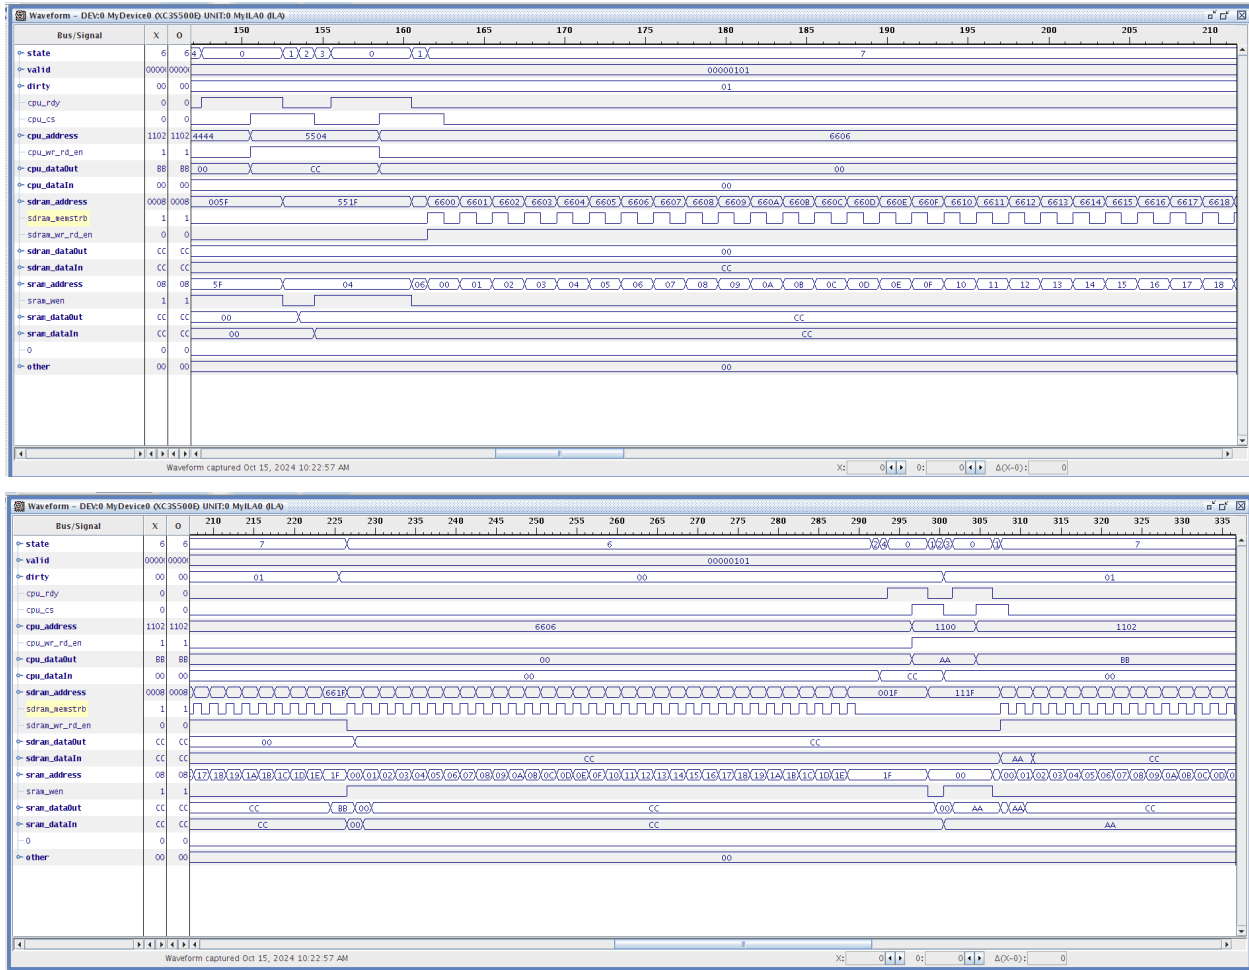


Figure 14. Cache operation (state 7) waveform

This waveform above represents when dirty bit = 1. Looking at x= 155, we can see that the cache controller is in an idle state, and the cache is waiting for a signal from the CPU to initiate memory operation. Once the signal is received, FSM will move on to state 1 (setting cpu_rdy to 0), where the cache controller checks whether it is a hit or a miss. State 1 concludes that it is a miss. Since it is a miss, it checks to see if the block is dirty. In this case, the dirty bit is 1 (the lsb of dirty is 1) and it will move onto state 7. In state 7, the cache controller writes the dirty block from the SRAM back to the SDRAM. In this case, it is the block at index 000. This state is responsible for writing back a cache block to the main memory (SDRAM) when a cache miss occurs and the cache block is dirty. It will write all 32 bytes starting with offset 0x00 and incrementing by one until 0x1F each clock cycle. This will ensure data consistency between the cache and the SDRAM. From the waveform, the SRAM loaded the value 0xCC to the SDRAM. Since the SDRAM is being written to, the sdram_wr_rd_en is set to 1. The new memory block then needs to be loaded from SDRAM into the SRAM after the write-back operation. This is done by going to state 6. The controller initiates the process of fetching the new block of data

from SDRAM and storing it in SRAM. This block will be written into the cache replacing the older cache line. The data will now perform its intended purpose by going to state 2 and determining whether the CPU is reading from or writing to the cache. It transitions to state 4 indicating that data the cache reads data to the CPU After it transitions to state 0.

Qualitative Results

N	Cache performance parameter	Time in nS
1	Hit / Miss determination time	10 ns
2	Data access time	10 ns
3	Block replacement time	650 ns
4	Hit time (Case 1 and 2)	30 ns
5	Miss penalty for Case 3 (when D-bit = 0)	Hit/Miss determination time + Block replacement time + hit time = 10ns + 650ns + 30ns = 690 ns
6	Miss penalty for Case 4 (when D-bit = 1)	Hit/Miss determination time + 2*Block replacement time + hit time = 10ns + 2*650ns + 30ns = 1340 ns

Figure 15. Cache performance parameters table

Conclusion

In this design project, the behavior of a cache controller was explored by implementing VHDL code using the Xilinx ISE CAD. The results of the cache controllers performance and its behaviour is detailed in Figure 15. These results agree with the theoretical principles of implementing cache controllers and their interactions with other system components.

During the project, we encountered numerous challenges, many of which were related to Chipscope errors, largely caused by human error of implementing the designed FSM in VHDL to correctly transition between states for each case. We also experienced difficulties with understanding how to use the ICON and ILA to trigger the CPU so it can send requests to the cache. Overall, the project provided a deep understanding on how the cache is used and greater exposure to using VHDL in designing digital systems.

References

1. COE758 Digital Systems Engineering Project #1 - Memory Hierarchy: Cache Controller, TMU, Toronto, ON, Canada, 2023. Accessed: Nov. 5, 2023. [Online].
[https://www.ecb.torontomu.ca/~lkirisch/ele758/labs/Cache%20Project\[12-09-10\].pdf](https://www.ecb.torontomu.ca/~lkirisch/ele758/labs/Cache%20Project[12-09-10].pdf)

Appendix

CACHE_CONTROLLER

```
32 entity cache_controller is
33   Port ( CLK_SRC      : in  std_logic;
34         CPU_ADDR      : out std_logic_vector(15 downto 0);
35         CPU_WR_RD      : out std_logic;
36         CS             : out std_logic;
37         CPU_DOUT       : out std_logic_vector(7 downto 0);
38         CPU_DIN        : out std_logic_vector(7 downto 0);
39         RDY            : out std_logic;
40         SDRAM_MEM_ADDR : out std_logic_vector(15 downto 0);
41         SDRAM_MEM_WR_RD : out std_logic;
42         MEMSTRB        : out std_logic;
43         FSM_STATE      : out std_logic_vector(2 downto 0);
44         SDRAM_DIN      : out std_logic_vector(7 downto 0);
45         SDRAM_DOUT     : out std_logic_vector(7 downto 0);
46         SRAM_DIN       : out std_logic_vector(7 downto 0);
47         SRAM_DOUT      : out std_logic_vector(7 downto 0);
48         SRAM_ADDR      : out std_logic_vector(7 downto 0));
49 end cache_controller;
50
51 architecture Behavioral of cache_controller is
52   -- table to store dirty, valid and cache tags
53   type memory is array (7 downto 0) of std_logic_vector(7 downto 0);
54   signal tag_reg : memory := (others => (others => ('0')));
55   signal valid   : std_logic_vector(7 downto 0) := "00000000";
56   signal dirty   : std_logic_vector(7 downto 0) := "00000000";
57
58   -- address word register fields
59   signal tag      : std_logic_vector(7 downto 0);
60   signal index    : std_logic_vector(2 downto 0);
61   signal offset   : std_logic_vector(4 downto 0);
62
63   -- sram signals
64   signal sram_wen      : std_logic_vector(0 downto 0);
65   signal sram_dataIn   : std_logic_vector(7 downto 0);
66   signal sram_dataOut  : std_logic_vector(7 downto 0);
67   signal sram_address  : std_logic_vector(7 downto 0);
68
69   -- sdram signals
70   signal cache_tag      : std_logic_vector(7 downto 0);
71   signal sdram_address  : std_logic_vector(15 downto 0);
72   signal sdram_wr_rd_en : std_logic;
73   signal sdram_memstrb  : std_logic;
74   signal sdram_dataIn   : std_logic_vector(7 downto 0);
75   signal sdram_dataOut  : std_logic_vector(7 downto 0);
76   signal counter        : integer;
77   signal mm_offset      : integer := 0;
78
79   -- cpu signals
80   signal cpu_rdy        : std_logic;
81   signal cpu_address    : std_logic_vector(15 downto 0);
82   signal cpu_wr_rd_en   : std_logic;
83   signal cpu_cs         : std_logic;
84   signal cpu_dataIn     : std_logic_vector(7 downto 0);
85   signal cpu_dataOut    : std_logic_vector(7 downto 0);
86   signal cpu_rst        : std_logic;
87
88   -- fsm signals
89   signal state : std_logic_vector(2 downto 0);
90   type state_type is (s1, s2, s3, s4, s6, s7, s0);
91   signal yfsm : state_type;
92   signal present_state : state_type;
93
94   -- ila, icon, vio signals
95   signal control0 : std_logic_vector(35 downto 0);
96   signal control1 : std_logic_vector(35 downto 0);
97   signal trig     : std_logic_vector(7 downto 0);
98   signal ila_data : std_logic_vector(119 downto 0);
```

```

100 -- SDRAM controller component
101 component sdram_controller
102     Port (clk : in std_logic;
103           addr : in std_logic_vector(15 downto 0);
104           memstrb : in std_logic;
105           wr_rd : in std_logic;
106           din : in std_logic_vector(7 downto 0);
107           dout : out std_logic_vector(7 downto 0));
108 end component;
109
110 component bram
111     port (clka : IN STD_LOGIC;
112           wea : IN STD_LOGIC_VECTOR(0 DOWNTO 0);
113           addra : IN STD_LOGIC_VECTOR(7 DOWNTO 0);
114           dina : IN STD_LOGIC_VECTOR(7 DOWNTO 0);
115           douta : OUT STD_LOGIC_VECTOR(7 DOWNTO 0));
116 end component;
117
118 -- CPU component
119 component CPU_gen
120     Port (clk : in STD_LOGIC;
121           rst : in STD_LOGIC;
122           trig : in STD_LOGIC;
123           Address : out STD_LOGIC_VECTOR(15 downto 0);
124           wr_rd : out STD_LOGIC;
125           cs : out STD_LOGIC;
126           DOut : out STD_LOGIC_VECTOR(7 downto 0));
127 end component;
128
129 component icon
130     port (CONTROL0 : INOUT STD_LOGIC_VECTOR(35 DOWNTO 0));
131 end component;
132

```

```

125
126     component icon
127     port (CONTROL0 : INOUT STD_LOGIC_VECTOR(35 DOWNTO 0));
128 end component;
129
130 -- ila component
131 component ila
132     port (CONTROL : INOUT STD_LOGIC_VECTOR(35 DOWNTO 0);
133           CLK : IN STD_LOGIC;
134           DATA : IN STD_LOGIC_VECTOR(119 DOWNTO 0);
135           TRIG0 : IN STD_LOGIC_VECTOR(7 DOWNTO 0));
136 end component;
137
138 begin
139
140 -- SDRAM instantiation
141 big_sdram: sdram_controller port map(
142     clk => CLK_SRC,
143     addr => sdram_address,
144     memstrb => sdram_memstrb,
145     wr_rd => sdram_wr_rd_en,
146     din => sdram_dataIn,
147     dout => sdram_dataOut);
148
149 -- SRAM instantiation
150 sram : bram port map (
151     clka => CLK_SRC,
152     wea => sram_wen,
153     addra => sram_address,
154     dina => sram_dataIn,
155     douta => sram_dataOut);
156
157 -- CPU instantiation
158 big_cpu: cpu_gen port map(
159     clk => CLK_SRC,
160     rst => '0',
161     trig => cpu_rdy,
162     Address => cpu_address,
163     wr_rd => cpu_wr_rd_en,
164     cs => cpu_cs,
165     DOut => cpu_dataOut);
166
167 -- icon instantiation
168 sys_icon : icon port map (CONTROL0 => control0);
169
170 -- ila instantiation
171 sys_ila : ila port map (
172     CONTROL => control0,
173     CLK => CLK_SRC,
174     DATA => ila_data,
175     TRIG0 => trig);

```

```

177 -- Next state generation
178 process(CLK_SRC)
179 begin
180   if (CLK_SRC'event and CLK_SRC='1') then
181     case yfsm is
182       when s1 => -- s1: checks for hit
183         cpu_rdy <= '0'; -- cache controller is not idle
184         state<="001";
185         tag <= cpu_address(15 downto 8); -- get tag from cpu address
186         index <= cpu_address(7 downto 5); -- get index from cpu address
187         offset <= cpu_address(4 downto 0); -- get offset from cpu address
188         sdram_address(15 downto 5) <= cpu_address(15 downto 5); -- get the tag and index from the cpu address (tag & index)
189         sdram_address(7 downto 0) <= cpu_address(7 downto 0); -- get the index and offset from cpu address (index & offset)
190         sram_wen <= "0"; -- not writing to sram
191         -- check that tag is in tag table and the valid is 1
192         if (valid(to_integer(unsigned(index))) = '1') and (tag = tag_reg(to_integer(unsigned(index)))) then
193           yfsm <= s2; -- found a hit so look for read/write operation
194         else
195           -- check dirty bit and valid
196           if (dirty(to_integer(unsigned(index))) = '1') and (valid(to_integer(unsigned(index))) = '1') then
197             yfsm<=s7; -- load sram block to sdram
198           else
199             yfsm<=s6; -- write new memory block to cache
200           end if;
201         end if;
202       when s2 => --s2: found a HIT
203         state<="010";
204         if (cpu_wr_rd_en='1') then
205           yfsm <= s3; -- write from cpu to cache
206         else
207           yfsm <= s4; -- read from cache to cpu
208         end if;
209       when s3 =>
210         state <= "011";
211         sram_wen <= "1"; -- writing to sram
212         valid(to_integer(unsigned(index))) <= '1'; -- set valid bit
213         dirty(to_integer(unsigned(index))) <= '1'; -- set dirty bit
214         sram_dataIn <= cpu_dataOut; -- loading the cpu data to the cache block
215         cpu_dataIn <= "00000000"; -- doesn't matter
216         yfsm<=s0; -- goes to idle
217       when s4 => -- s4: reading from cpu to cache
218         state <= "100";
219         cpu_dataIn <= sram_dataOut; -- output cache value to cpu
220         yfsm<=s0; -- go to idle
221       when s6=> -- s6: writing new memory block to cache
222         state<="110";
223         if (counter = 64) then -- takes 32 cycles to write to memory
224           counter<=0;
225           valid(to_integer(unsigned(index))) <= '1'; -- set valid to 1
226           tag_reg(to_integer(unsigned(index))) <= tag; -- get tag from tag register
227           mm_offset<=0;
228           yfsm<=s2; -- check for read or
229         else
230           if (counter mod 2 = 1) then -- only access memory on rising edges
231             sdram_memstrib<='0';
232           else
233             sdram_address <= cache_tag & index & std_logic_vector(to_unsigned(mm_offset, 5)); -- (tag & index & offset)
234             sdram_wr_rd_en <= '0'; -- not writing to sdram
235             sdram_memstrib <= '1';
236             sram_address <= index & std_logic_vector(to_unsigned(mm_offset, 5)); -- (index & offset)
237             sram_dataIn <= sdram_dataOut; -- writing sdram block to sram
238             sram_wen <= "1"; -- enable sram for writing
239             mm_offset <= mm_offset+1; -- increment to next byte
240           end if;
241           counter <= counter + 1; -- increment counter
242         end if;
243       when s7=> -- s7: write cache block to memory
244         state<="111";
245         if (counter = 64) then -- takes 32 cycles to read from memory
246           dirty(to_integer(unsigned(index))) <= '0'; -- set dirty to 0
247           tag_reg(to_integer(unsigned(index))) <= tag; -- the tag in tag table
248           counter<=0; --reset counters
249           mm_offset<=0;
250           yfsm<=s6; -- load new block from sdram to sram
251         else
252           if (counter mod 2 = 1) then -- writes only on rising edge
253             sdram_memstrib<='0';
254           else
255             sdram_address <= tag & index & std_logic_vector(to_unsigned(mm_offset, 5)); -- write each byte in the block (tag&index&offset)
256             sdram_wr_rd_en <= '1'; -- writing to sdram
257             sram_address <= index & std_logic_vector(to_unsigned(mm_offset, 5)); -- (index&offset)
258             sram_wen <= "0"; -- not writing to cache
259             --sram_dataIn <= sdram_dataOut; --
260             sdram_dataIn <= sram_dataOut; -- writing cache data to memory
261             sdram_memstrib<='1';
262             mm_offset <= mm_offset+1; -- increment to next byte
263           end if;
264           counter <= counter + 1; -- increment counter
265         end if;
266       when s0 => -- s0: cache idle
267         state <= "000";
268         cpu_rdy <= '1'; -- tell cpu cache is idle
269         if (cpu_cs = '1') then -- cpu sends signal
270           yfsm <= s1; -- check for hit/miss
271         end if;
272       end case;
273     end if;
274   end process;
275

```

```

276 -- store data in signals to ports
277 FSM_STATE <= state;
278 CPU_ADDR <= cpu_address;
279 CPU_WR_RD <= cpu_wr_rd_en;
280 --CS <= cpu_cs;
281 CPU_DOUT <= cpu_dataOut;
282 CPU_DIN <= cpu_dataIn;
283 RDY <= cpu_rdy;
284 SDRAM_MEM_ADDR <= sdram_address;
285 SDRAM_MEM_WR_RD <= sdram_wr_rd_en;
286 MEMSTRB <= sdram_memstrb;
287 SDRAM_DIN <= sdram_dataIn;
288 SDRAM_DOUT <= sdram_dataOut;
289 SRAM_DIN <= sram_dataIn;
290 SRAM_DOUT <= sram_dataOut;
291 SRAM_ADDR <= sram_address;
292
293 -- map ports to ila data
294 ila_data(15 downto 0) <= cpu_address;
295 ila_data(16) <= cpu_wr_rd_en;
296 ila_data(17) <= cpu_cs;
297 ila_data(18) <= '0';
298 ila_data(19) <= cpu_rdy;
299 ila_data(27 downto 20) <= cpu_dataOut;
300 ila_data(35 downto 28) <= cpu_dataIn;
301
302 ila_data(36) <= sdram_memstrb;
303 ila_data(52 downto 37) <= sdram_address;
304 ila_data(53) <= sdram_wr_rd_en;
305 ila_data(61 downto 54) <= sdram_dataIn;
306 ila_data(69 downto 62) <= sdram_dataOut;
307
308 ila_data(77 downto 70) <= sram_dataIn;
309 ila_data(85 downto 78) <= sram_dataOut;
310 ila_data(93 downto 86) <= sram_address;
311 ila_data(94 downto 94) <= sram_wen;
312 ila_data(97 downto 95) <= state;
313
314 ila_data(105 downto 98) <= valid;
315 ila_data(113 downto 106) <= dirty;
316 ila_data(119 downto 114) <= (others=>'0');
317
318 end Behavioral;

```

CPU_GEN

```

27 entity CPU_gen is
28   Port (
29     clk      : in  STD_LOGIC;
30     rst      : in  STD_LOGIC;
31     trig     : in  STD_LOGIC;
32     -- Interface to the Cache Controller.
33     Address  : out STD_LOGIC_VECTOR (15 downto 0);
34     wr_rd    : out STD_LOGIC;
35     cs       : out STD_LOGIC;
36     DOut     : out STD_LOGIC_VECTOR (7 downto 0)
37   );
38 end CPU_gen;
39
40 architecture Behavioral of CPU_gen is
41
42   -- Pattern storage and control.
43   signal patOut : std_logic_vector(24 downto 0);
44   signal patCtrl : std_logic_vector(2 downto 0) := "111";
45   signal updPat : std_logic;
46
47   -- Main control.
48   signal stl : std_logic_vector(2 downto 0) := "000";
49   signal stlN : std_logic_vector(2 downto 0);
50
51   -- Rising edge detection.
52   signal rReg1, rReg2 : std_logic;
53   signal trig_r : std_logic;
54
55 begin
56
57   -----
58   -- Main control FSM.
59   -----
60
61   -- State storage.
62   process(clk, rst, stlN)
63   begin
64     if(rst = '1')then
65       stl <= "000";
66     else
67       if(clk'event and clk = '1')then
68         stl <= stlN;
69       end if;
70     end if;
71   end process;

```

```

73  -- Next state generation.
74  process(st1, trig_r)
75  begin
76      if(st1 = "000")then
77          if(trig_r = '1')then
78              st1N <= "001";
79          else
80              st1N <= "000";
81          end if;
82      elsif(st1 = "001")then
83          st1N <= "010";
84      elsif(st1 = "010")then
85          st1N <= "011";
86      elsif(st1 = "011")then
87          st1N <= "100";
88      elsif(st1 = "100")then
89          st1N <= "101";
90      elsif(st1 = "101")then
91          st1N <= "000";
92      else
93          st1N <= "000";
94      end if;
95  end process;

```

```

97  -- Output generation.
98  process(st1)
99  begin
100      if(st1 = "000")then
101          updPat <= '0';
102          cs <= '0';
103      elsif(st1 = "001")then
104          updPat <= '1';
105          cs <= '0';
106      elsif(st1 = "010")then
107          updPat <= '0';
108          cs <= '1';
109      elsif(st1 = "011")then
110          updPat <= '0';
111          cs <= '1';
112      elsif(st1 = "100")then
113          updPat <= '0';
114          cs <= '1';
115      elsif(st1 = "101")then
116          updPat <= '0';
117          cs <= '1';
118      else
119          end if;
120  end process;

```

```

121  -----
122  -- Pattern generator and control circuit.
123

```

```

126  -- Generator control circuit.
127  process(clk, rst, updPat, patCtrl)
128  begin
129      if(rst = '1')then
130          patCtrl <= "111";
131      else
132          if(clk'event and clk = '1')then
133              if(updPat = '1')then
134                  patCtrl <= patCtrl + "001";
135              else
136                  patCtrl <= patCtrl;
137              end if;
138          end if;
139      end if;
140  end process;
141
142  -- Pattern storage.
143  process(patCtrl)
144  begin
145      if(patCtrl = "000")then
146          patOut <= "0001000100000000101010101";
147      elsif(patCtrl = "001")then
148          patOut <= "00010001000000010101110111";
149      elsif(patCtrl = "010")then
150          patOut <= "00010001000000000000000000";

```

```

151     elsif(PatCtrl = "011")then
152         patOut <= "0001000100000001000000000000";
153     elsif(PatCtrl = "100")then
154         patOut <= "0011001101000110000000000000";
155     elsif(PatCtrl = "101")then
156         patOut <= "0100010001000100000000000000";
157     elsif(PatCtrl = "110")then
158         patOut <= "01010101000000100110011001";
159     else
160         patOut <= "0110011000000011000000000000";
161     end if;
162 end process;
163
164 -----
165 -- Rising edge detector.
166 -----
167
168 -- Register 1
169 process(clk, trig)
170 begin
171     if(clk'event and clk = '1')then
172         rReg1 <= trig;
173     end if;
174 end process;
175
176 -- Register 2
177 process(clk, rReg1)
178 begin
179     if(clk'event and clk = '1')then
180         rReg2 <= rReg1;
181     end if;
182 end process;
183
184 trig_r <= rReg1 and (not rReg2);
185
186 -----
187 -- Output connections.
188 -----
189
190 -- Output mapping:
191 -- Address [24 .. 9]
192 -- Data [8 .. 1]
193 -- Wr/Rd [0]
194
195 Address(15 downto 0) <= patOut(24 downto 9);
196 DOut(7 downto 0) <= patOut(8 downto 1);
197 wr_rd <= patOut(0);
198
199
200 end Behavioral;

```

ILA

```

20 ----- Begin Cut here for COMPONENT Declaration ----- COMP_TAG
21 component ila
22     PORT (
23         CONTROL : INOUT STD_LOGIC_VECTOR(35 DOWNTO 0);
24         CLK : IN STD_LOGIC;
25         DATA : IN STD_LOGIC_VECTOR(119 DOWNTO 0);
26         TRIG0 : IN STD_LOGIC_VECTOR(7 DOWNTO 0));
27 end component;
28
29
30 -- COMP_TAG_END ----- End COMPONENT Declaration -----
31 -- The following code must appear in the VHDL architecture
32 -- body. Substitute your own instance name and net names.
33 ----- Begin Cut here for INSTANTIATION Template ----- INST_TAG
34
35 your_instance_name : ila
36     port map (
37         CONTROL => CONTROL,
38         CLK => CLK,
39         DATA => DATA,
40         TRIG0 => TRIG0);

```


ICON

```
21 component icon
22   PORT (
23     CONTROL0 : INOUT STD_LOGIC_VECTOR(35 DOWNTO 0));
24
25 end component;
26
27 -- COMP_TAG_END ----- End COMPONENT Declaration -----
28 -- The following code must appear in the VHDL architecture
29 -- body. Substitute your own instance name and net names.
30 ----- Begin Cut here for INSTANTIATION Template ----- INST_TAG
31
32 your_instance_name : icon
33   port map (
34     CONTROL0 => CONTROL0);
35
36 -- INST_TAG_END ----- End INSTANTIATION Template -----
37
```

BRAM

```
50
51 ----- Begin Cut here for COMPONENT Declaration ----- COMP_TAG
52 COMPONENT bram
53   PORT (
54     clka : IN STD_LOGIC;
55     wea : IN STD_LOGIC_VECTOR(0 DOWNTO 0);
56     addra : IN STD_LOGIC_VECTOR(7 DOWNTO 0);
57     dina : IN STD_LOGIC_VECTOR(7 DOWNTO 0);
58     douta : OUT STD_LOGIC_VECTOR(7 DOWNTO 0)
59   );
60 END COMPONENT;
61 -- COMP_TAG_END ----- End COMPONENT Declaration -----
62
63 -- The following code must appear in the VHDL architecture
64 -- body. Substitute your own instance name and net names.
65 ----- Begin Cut here for INSTANTIATION Template ----- INST_TAG
66 your_instance_name : bram
67   PORT MAP (
68     clka => clka,
69     wea => wea,
70     addra => addra,
71     dina => dina,
72     douta => douta
73   );
74
75 -- INST_TAG_END ----- End INSTANTIATION Template -----
```

CACHE CONSTRAINTS

```
1 #Clk Signal
2 Net "CLK_SRC" loc = "C9" ;|
```